



Data modeling on Amazon DynamoDB

Aman Dhingra

DynamoDB Specialist SA, AWS

2021-09-29



Agenda

Amazon DynamoDB key concepts

NoSQL Workbench Demo

Hands-on lab

Quiz

Wrap-up

Application architecture and patterns have evolved

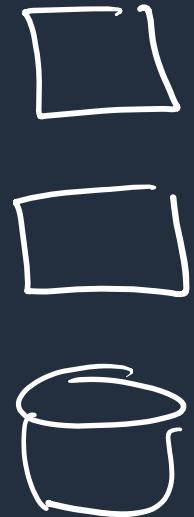
Mainframe



Client Server



Three Tier



Microservices



Characteristics of internet-scale apps



E-commerce



Media streaming



Social media



Online gaming



Shared economy

Users	1 million+
Data volume	TB, PB, EB
Locality	Global
Performance	Microsecond latency
Request rate	Millions per second
Access	Mobile, IoT, devices
Scale	Up and down
Economics	Pay as you go
Developer access	Instant API access

DynamoDB



Performance at scale

- Handles millions of requests per second
- Delivers single-digit-millisecond latency
- Automated global replication
- New advanced streaming with Amazon Kinesis Data Streams for DynamoDB



No servers to manage

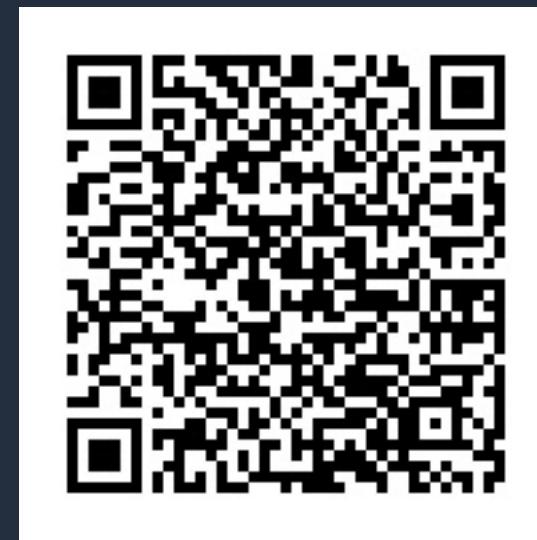
- Maintenance free
- Auto scaling
- On-demand capacity mode
- Change data capture for integration with AWS Lambda, Amazon Redshift, Amazon Elasticsearch Service
- AWS Glue Elastic Views for DynamoDB



Enterprise ready

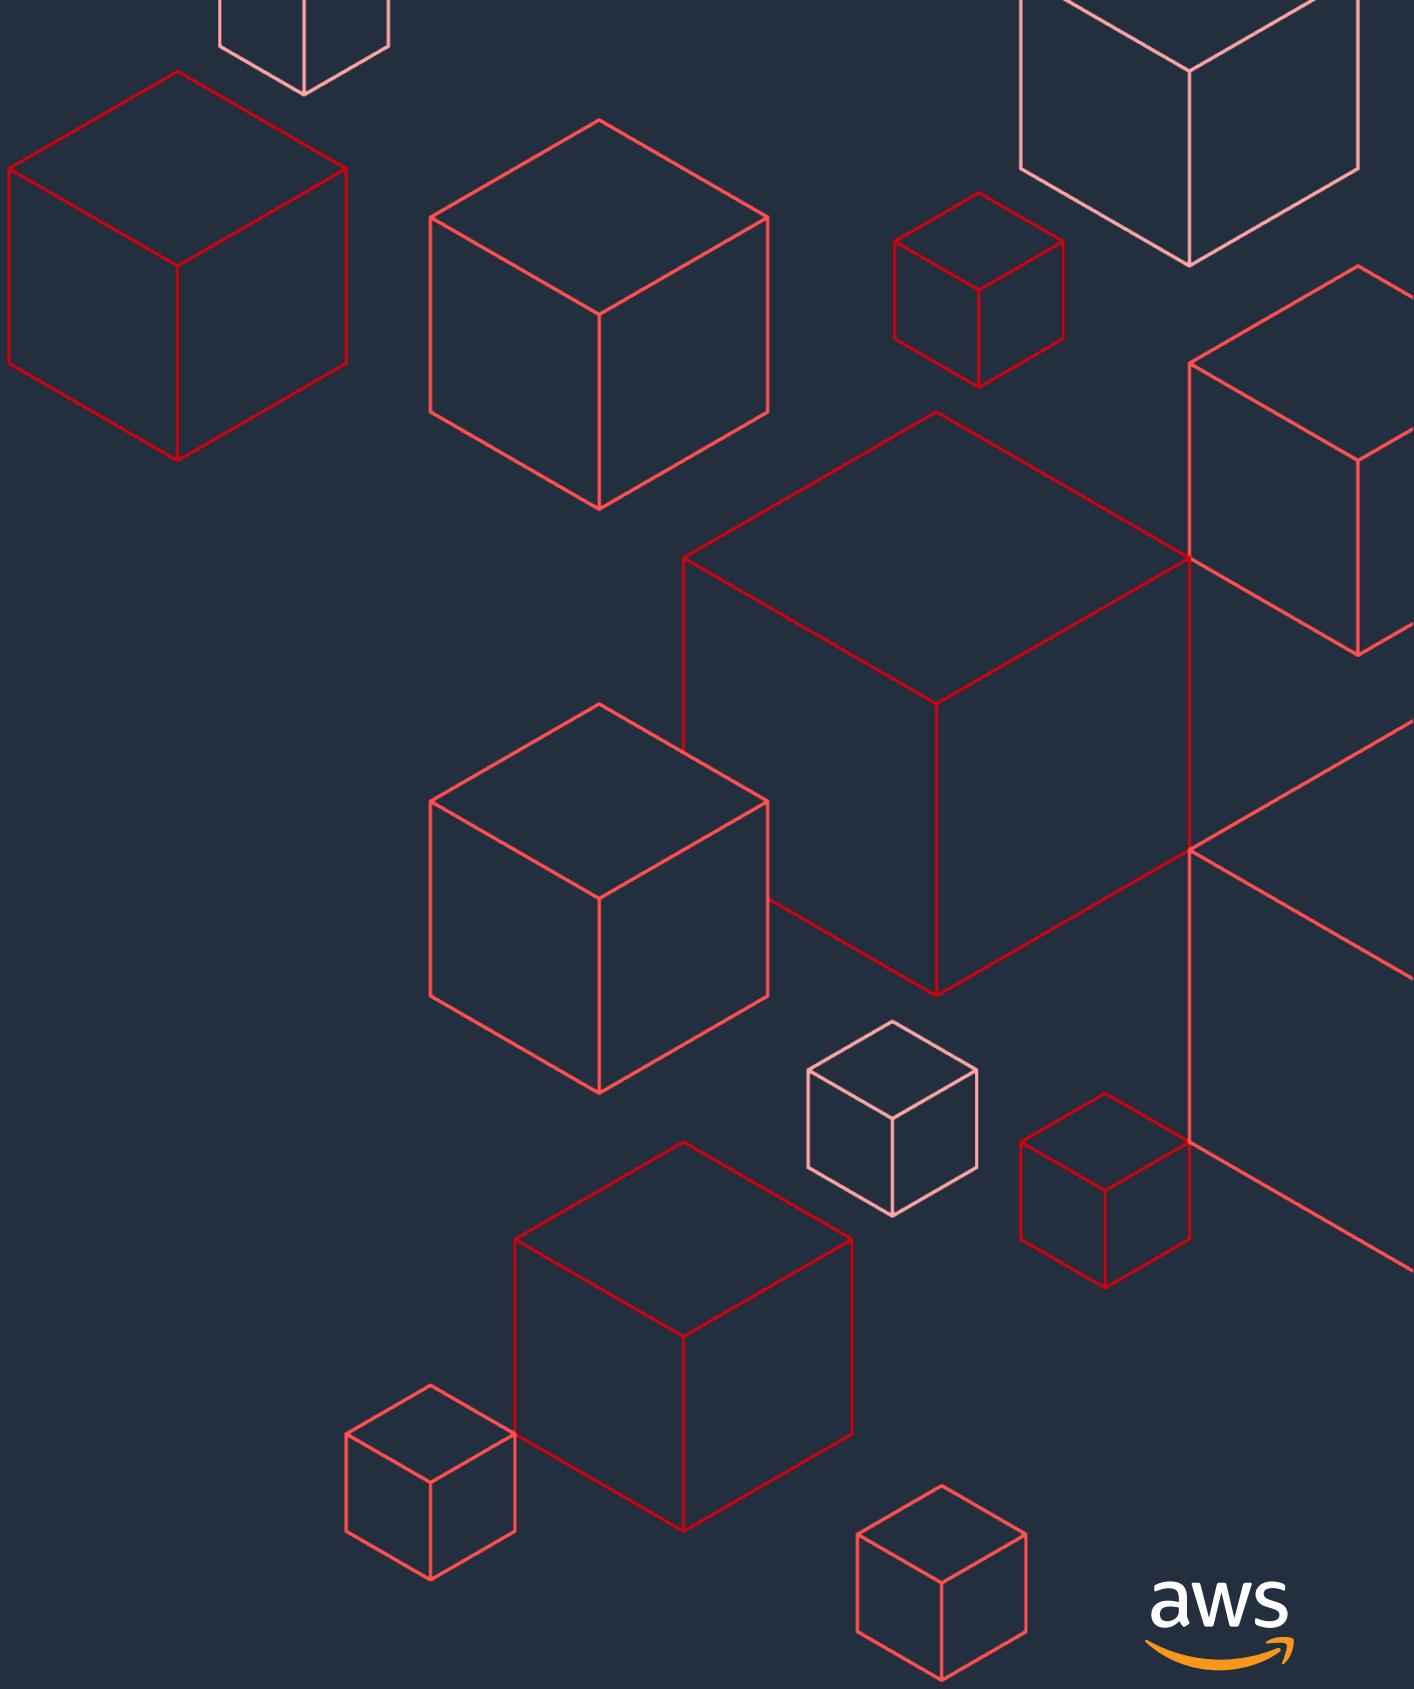
- ACID transactions
- Encryption at rest
- Continuous backups (PITR), and on-demand backup and restore
- NoSQL Workbench
- Export table data to S3
- PartiQL (a SQL-compatible query language) support

More about DynamoDB features



tinyurl.com/4dsa5jdn

Key concepts



Create table

Table details Info

DynamoDB is a schemaless database that requires only a table name and a primary key when you create the table.

Table name

This will be used to identify your table.

Between 3 and 255 characters, containing only letters, numbers, underscores (_), hyphens (-), and periods (.)



Partition key

The partition key is part of the table's primary key. It is a hash value that is used to retrieve items from your table and allocate data across hosts for scalability and availability.



1 to 255 characters and case sensitive.

Sort key - optional

You can use a sort key as the second part of a table's primary key. The sort key allows you to sort or search among all items sharing the same partition key.



1 to 255 characters and case sensitive.

Settings

Default settings

The fastest way to create your table. You can modify these settings now or after your table has been created.

Customize settings

Use these advanced features to make DynamoDB work better for your needs.

Default settings

Read/write capacity [Info](#)

Using provisioned capacity mode. Read and write capacity are set to 5 units each with auto scaling enabled.

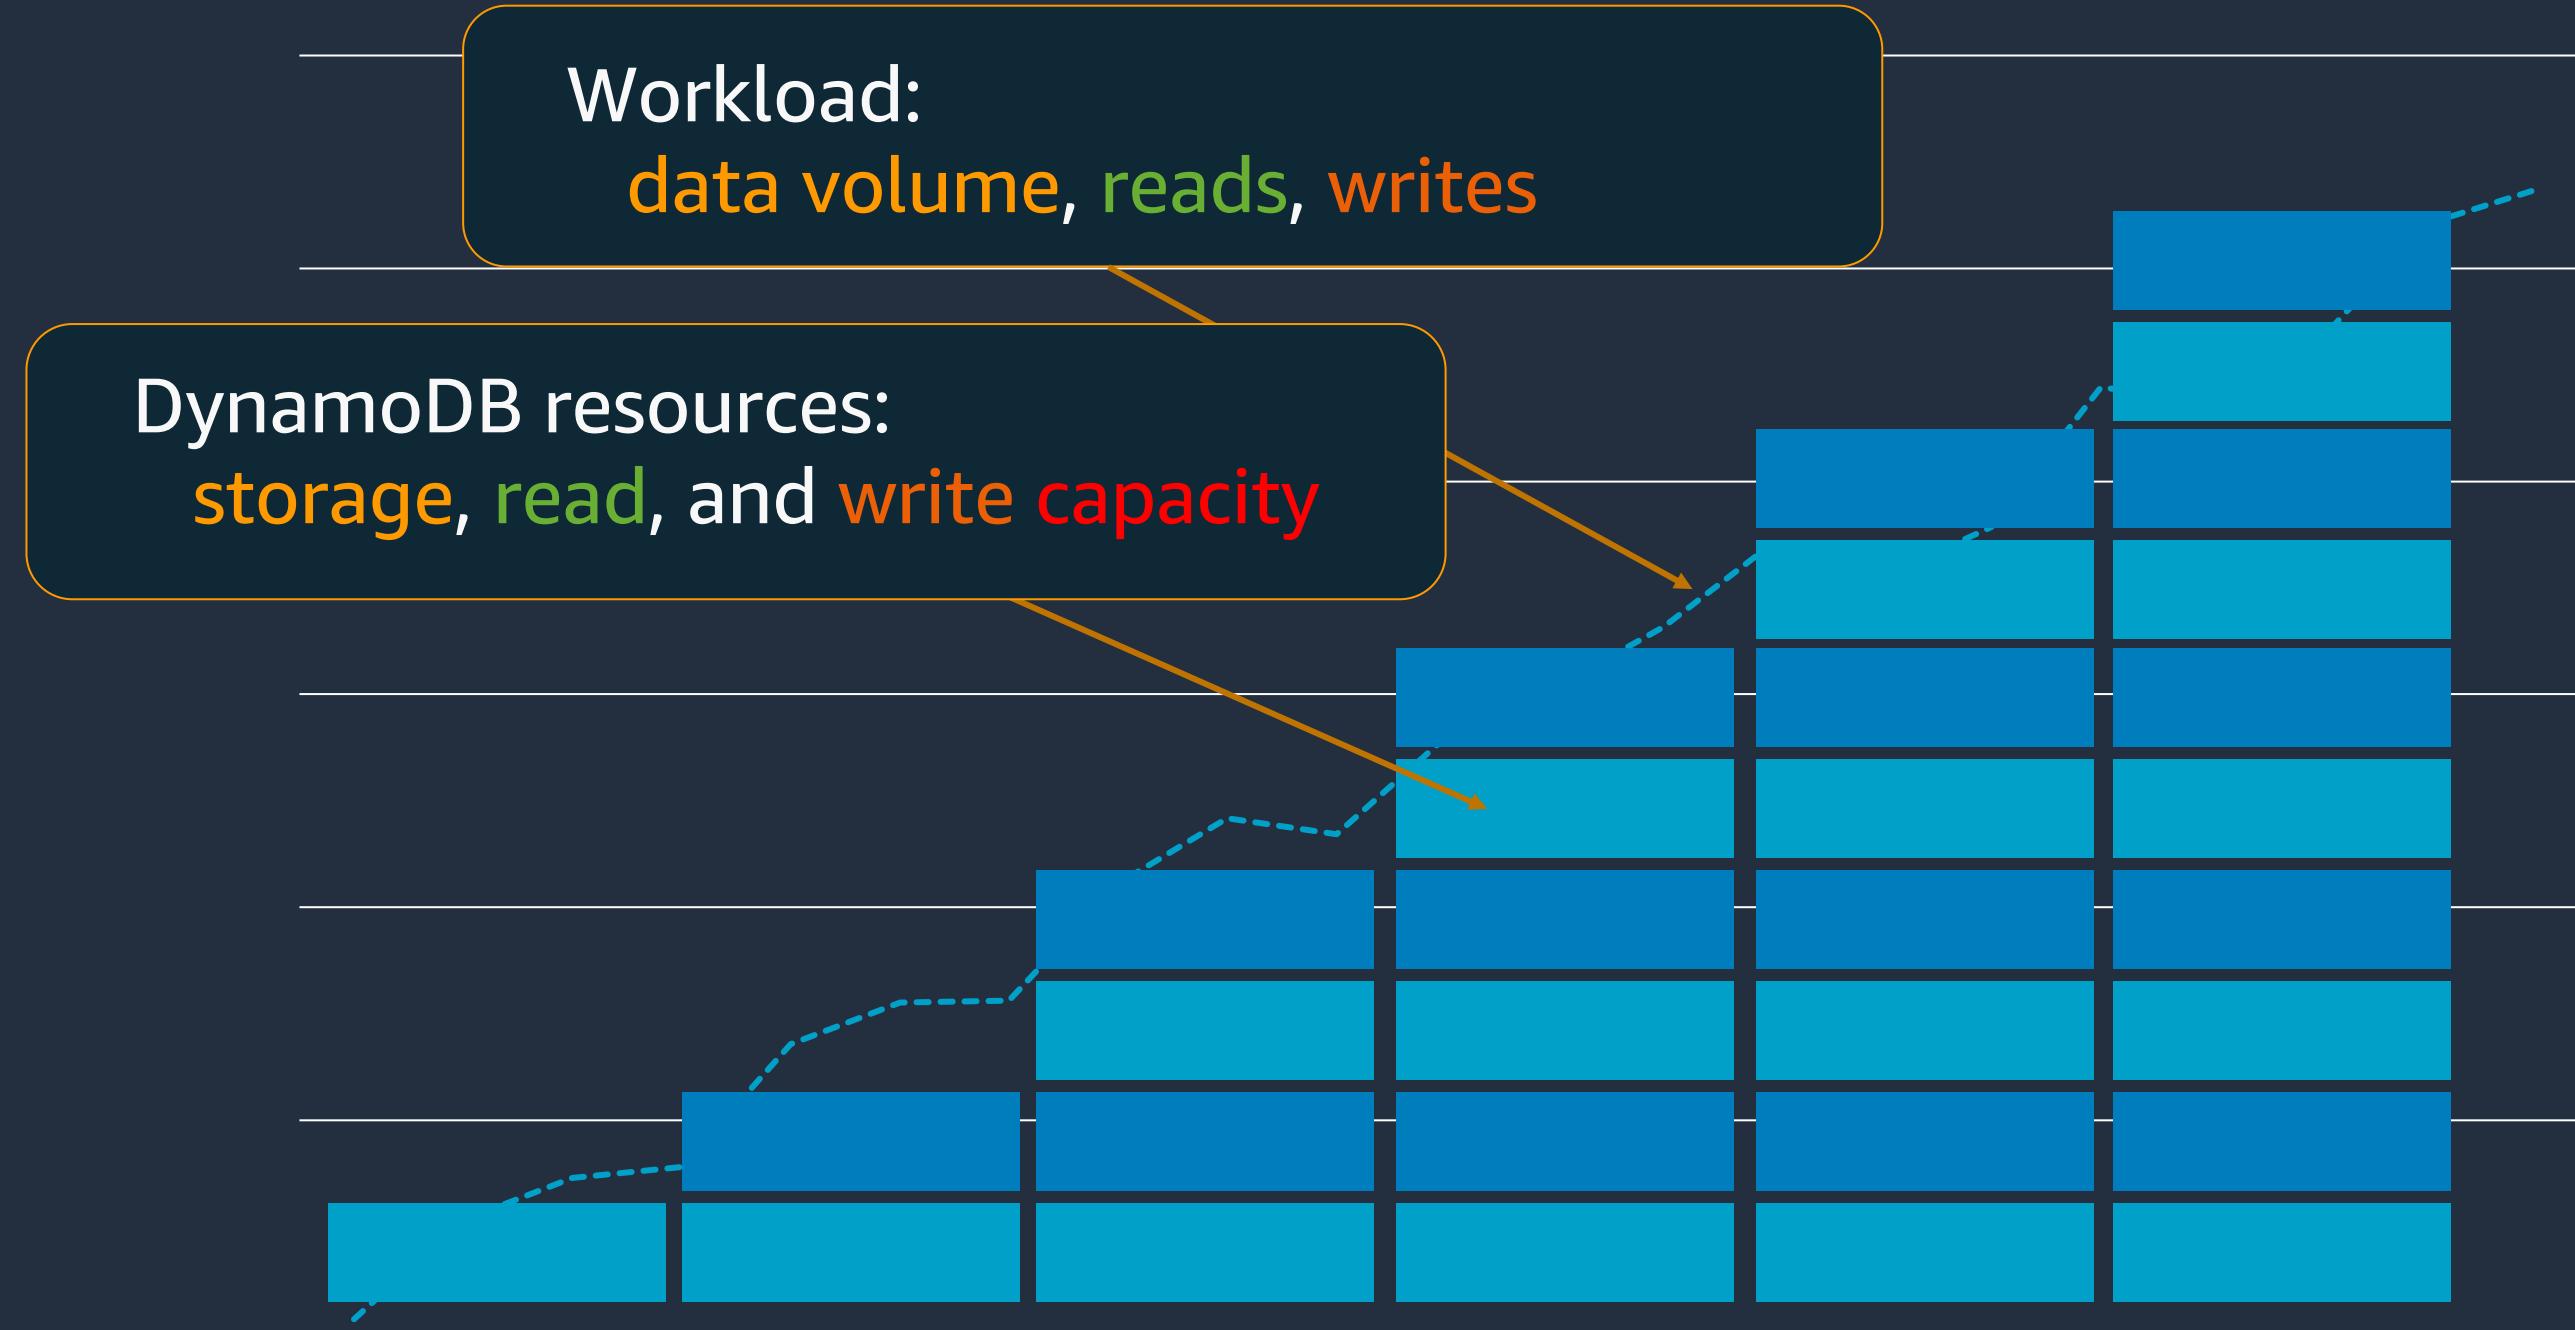
Secondary indexes [Info](#)

No secondary indexes have been created. Queries will be run by using the table's partition key and sort key only.

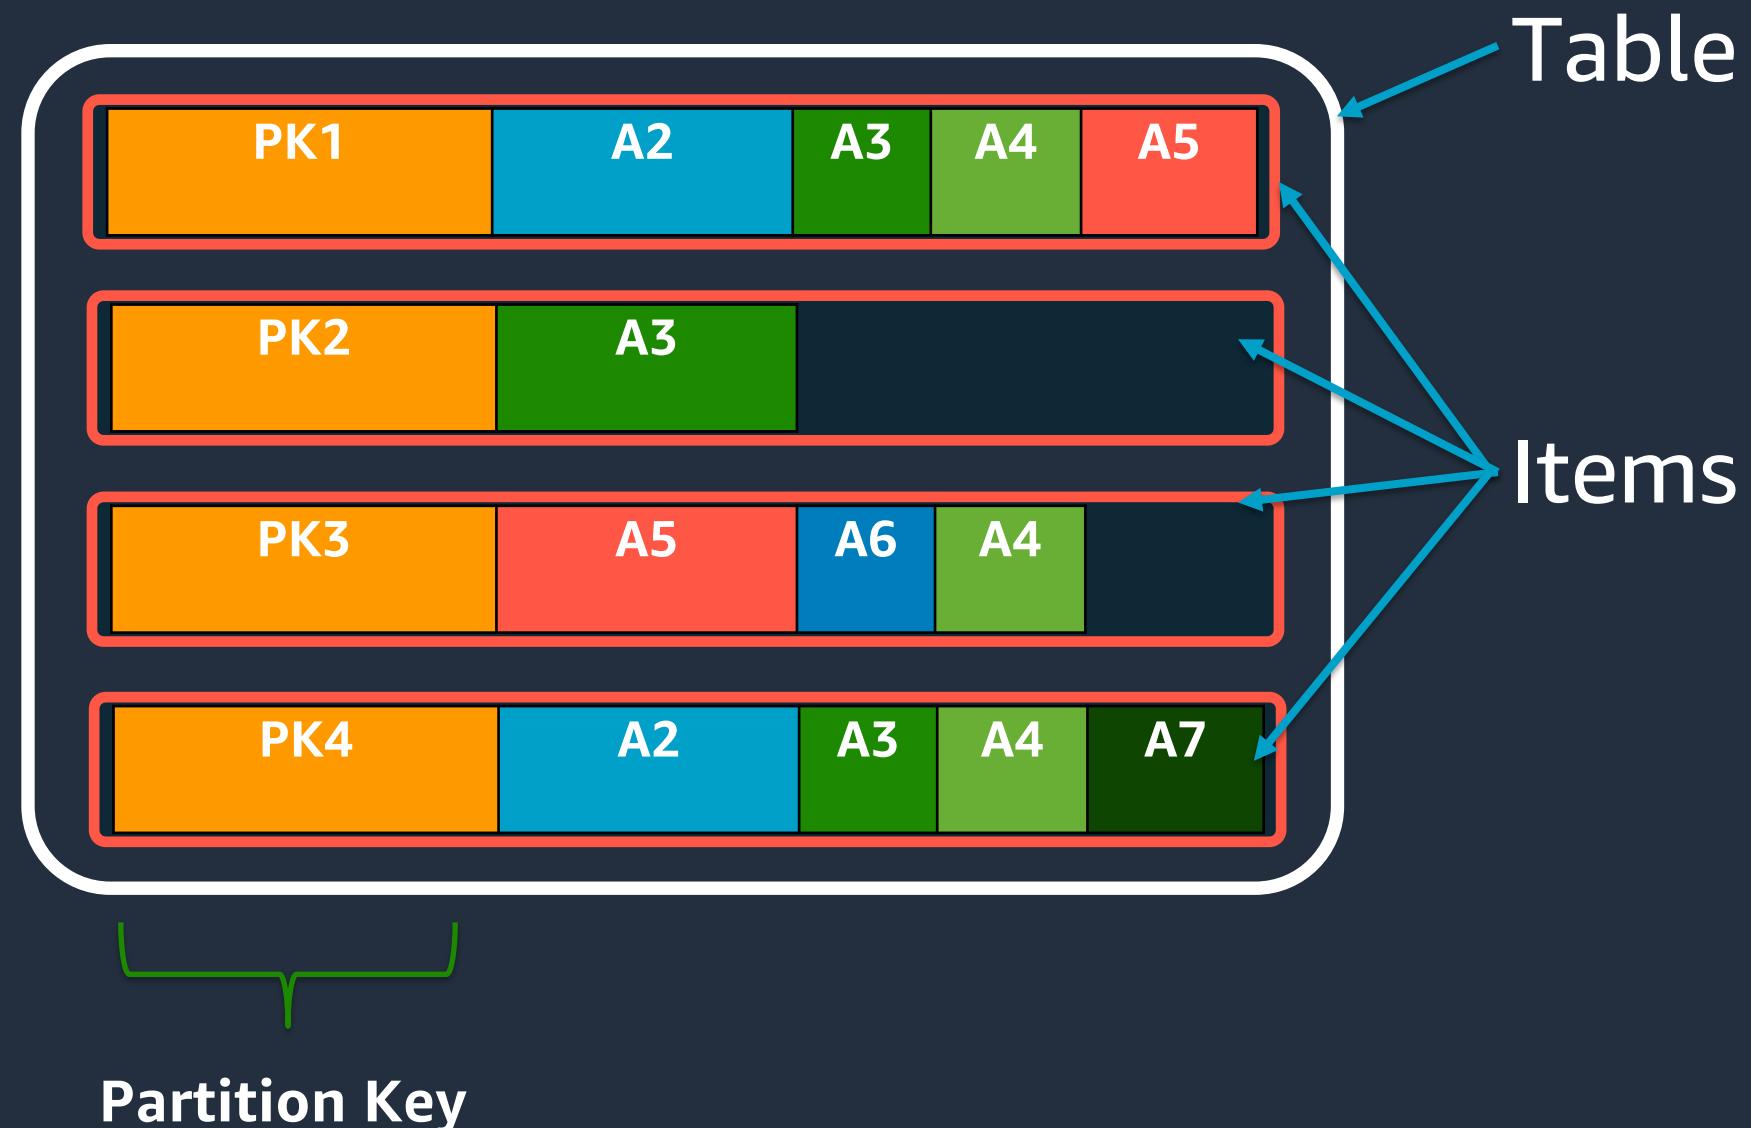
Key management for encryption at rest [Info](#)

Using the AWS owned customer master key. This key is managed by DynamoDB at no extra cost.

Horizontal scaling with DynamoDB

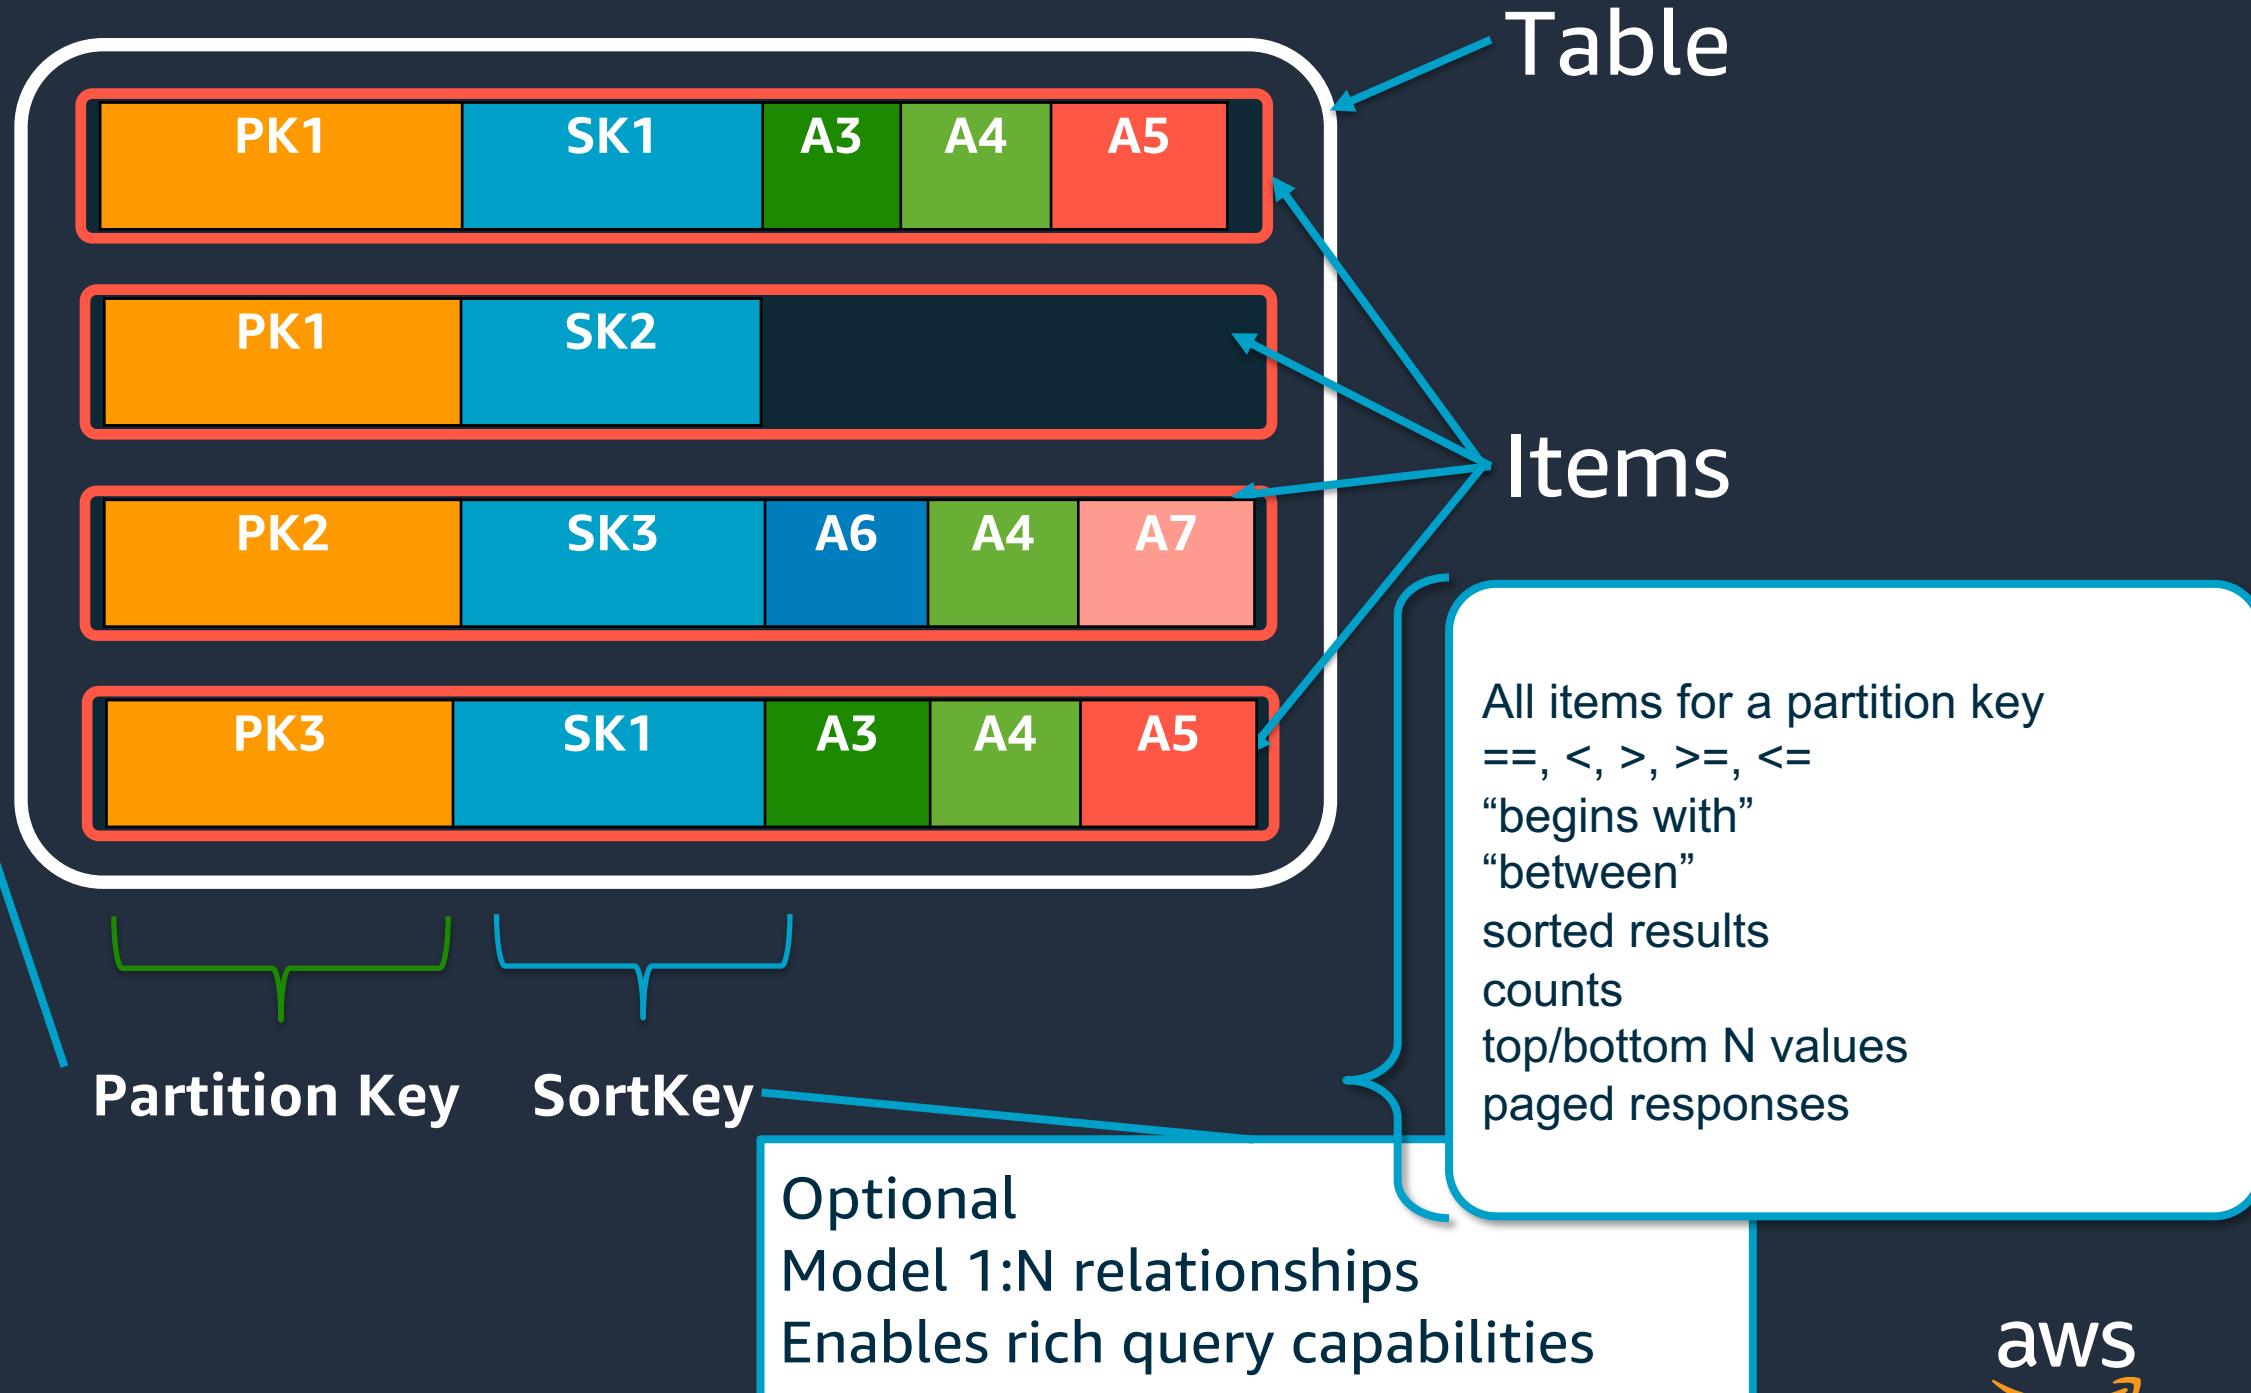


DynamoDB Table



DynamoDB Table

Mandatory
Key-value access pattern
Determines data
distribution



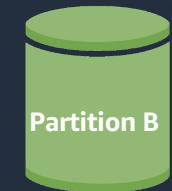
A view from “a different angle”



Table



Host 1



Host 2



Host 3

Availability Zone A



Host 4



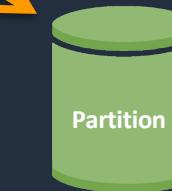
Host 5



Host 6



Host 7



Host 8



Host 9

Availability Zone B

Availability Zone C

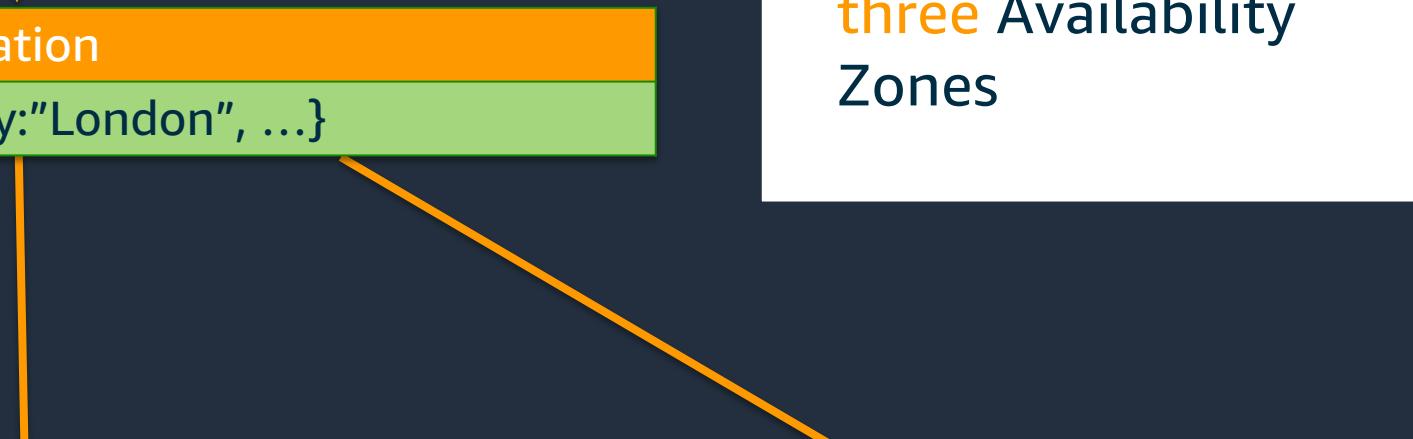
	CustID	Customer Information
	523422	{ name:"Alex", city:"London", ...}

Hash(523422) = 0xF355

Hash Value	CustID	Customer Information
0xF355	523422	{ name:"Alex", city:"London", ...}



Host 5



Three-way replication

- Data is always replicated to **three** Availability Zones
- The service runs in **three** Availability Zones

Local secondary index (LSI)

Table	{	A1 (PK)	A2 (SK)	A3	A4	A5
-------	---	------------	------------	----	----	----

Added at table creation time.

RCUs/WCUs are consumed from base Table.

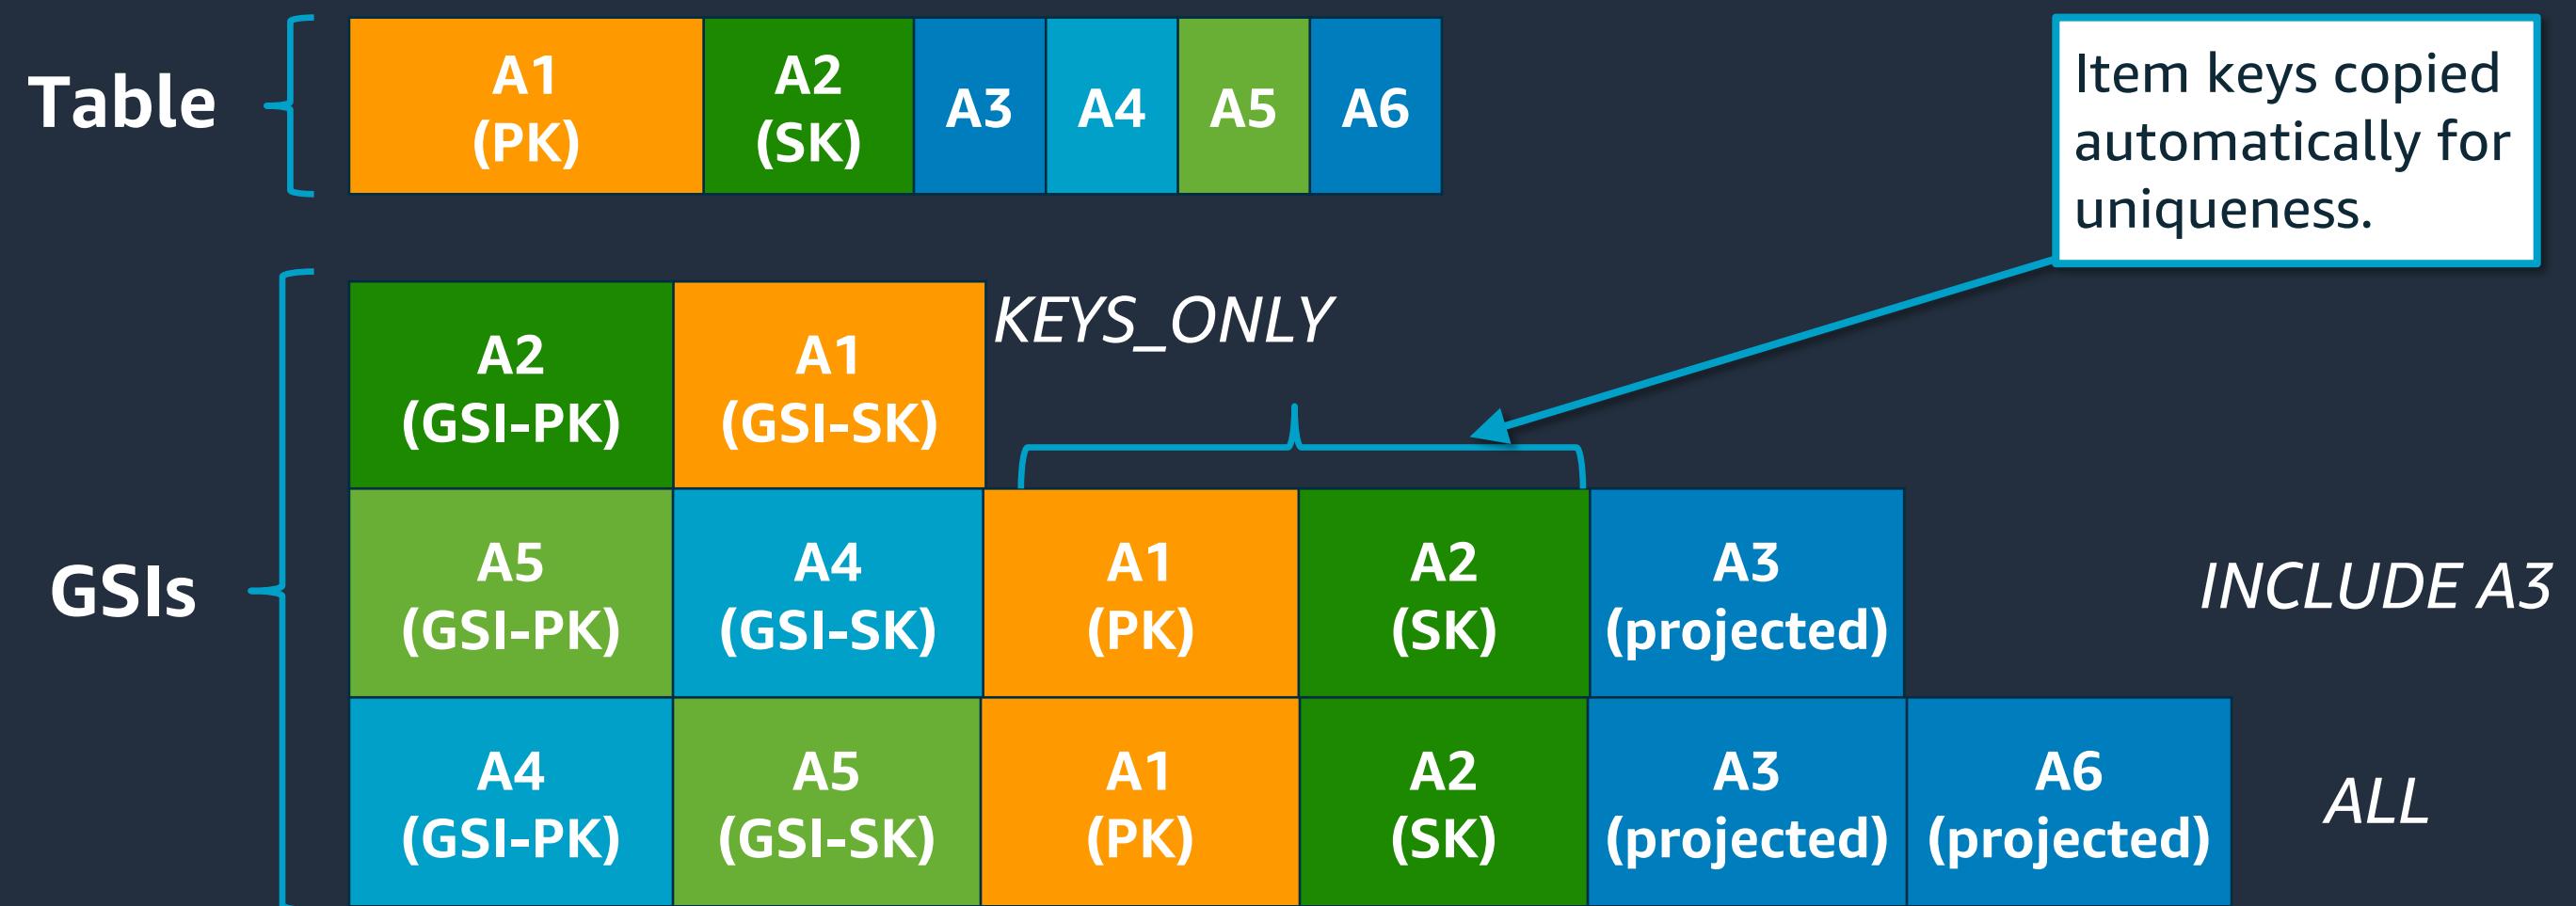
Up to 5 LSIs per table.

LSIs	{	A1 (PK)	A3 (LSI-SK)	A2 (SK)	KEYS_ONLY	Table's sort key is copied automatically.
		A1 (PK)	A4 (LSI-SK)	A2 (SK)	A3 (projected)	INCLUDE A3
		A1 (PK)	A5 (LSI-SK)	A2 (SK)	A3 (projected)	A4 (projected)

Global secondary index (GSI)

Up to 20 GSIs per table.

RCUs/WCUs provisioned separately for GSIs.



DynamoDB Types		Data Types
Scalars	Number	N Integer Float Timestamp
	String	S Timestamp Date (ISO 8601)
	Binary	B Blob
	Boolean	BOOL Boolean
Document	Null	NULL Null
	List	L List
	Map	M Map
Set	String Set	SS
	Number Set	NS Set
	Binary Set	BS

DynamoDB Operations

Working with specific Items

DynamoDB Operations	Expressions
GetItem	Projection Expression
BatchGetItem	Projection Expression
TransactReadItems	Projection Expression
PutItem	Condition Expression
BatchWriteItem	
UpdateItem	Condition Expression Update Expression
DeleteItem	Condition Expression
TransactWriteItems	Condition Expression

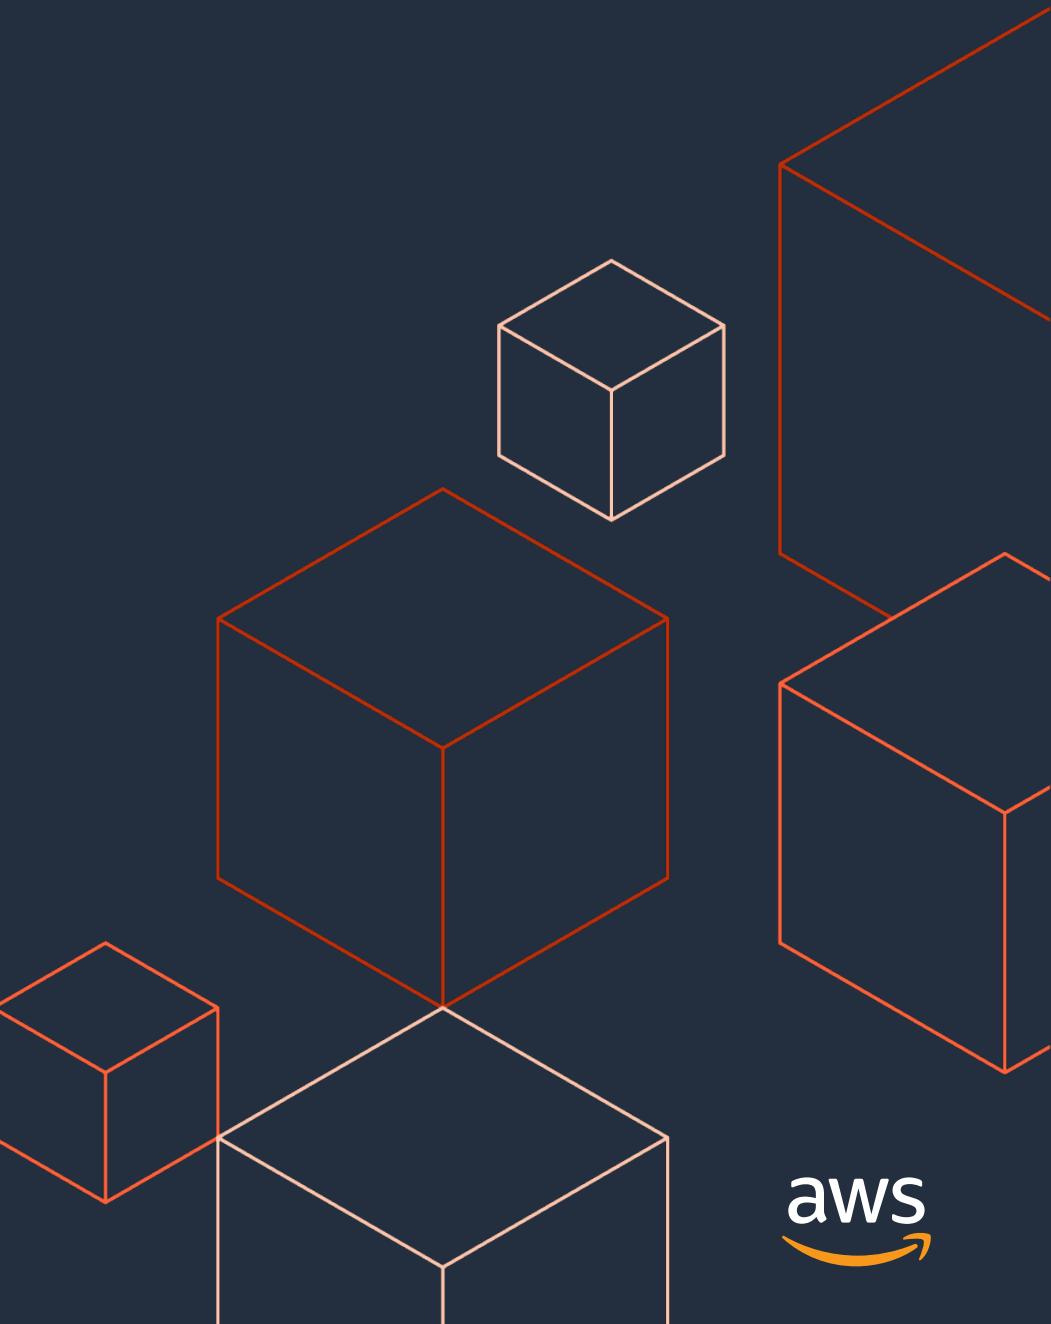
DynamoDB Operations		Expressions
Working with an Item collection	Query	Key Condition Expression
		Filter Expression Projection Expression
Working with the entire table	Scan	Filter Expression Projection Expression

PartiQL APIs

Operations for PartiQL:

- *ExecuteStatement*
 - Support for single or multiple item SELECT
 - Support for single item INSERT, UPDATE or DELETE
- *BatchExecuteStatement*
 - Support for a batch of single item SELECT OR batch of single item INSERT/DELETE of up to 25 items
- *ExecuteTransaction*
 - Support for a batch of single item SELECT OR batch of single item INSERT/UPDATE/DELETE/CONDITION CHECK of up to 25 items

NoSQL data modeling



All data is relational



IT monitoring



Social graph



Documents



Auth services

Tenets of DynamoDB data modeling

- Understand the use case
- Identify the access patterns
 - Read/write workloads
 - Query dimensions and aggregations
- Data modeling
 - Using NoSQL design patterns
- Review -> Repeat -> Review
- Nature of the data
 - OLTP / OLAP / full-text search
 - Relationships between the entities
 - What does concurrent access look like?
 - Time series data
 - Archiving needs, etc.

Tenets of DynamoDB data modeling

- Understand the use case
- Identify the access patterns
 - Read/write workloads
 - Query dimensions and aggregations
- Data modeling
 - Using NoSQL design patterns
- Review -> Repeat -> Review
- Source data analysis (write workload)
- Reading one item versus multiple items (read workload)
- Query aggregations and KPIs

Tenets of DynamoDB data modeling

- Understand the use case
- Identify the access patterns
 - Read/write workloads
 - Query dimensions and aggregations
- Data modeling
 - Using NoSQL design patterns
- Review -> Repeat -> Review
- 1:1, 1:n, m:n relationships
- **1 application = 1 table**
 - Avoid unnecessary fetches
 - Simplify access patterns
- Identify primary key
 - Partition key and Sort key
- Query dimensions using LSIs and GSIs

Tenets of DynamoDB data modeling

- Understand the use case
- Identify the access patterns
 - Read/write workloads
 - Query dimensions and aggregations
- Data modeling
 - Using NoSQL design patterns
- Review -> Repeat -> Review

Use case Discovery Phase

Online Shop

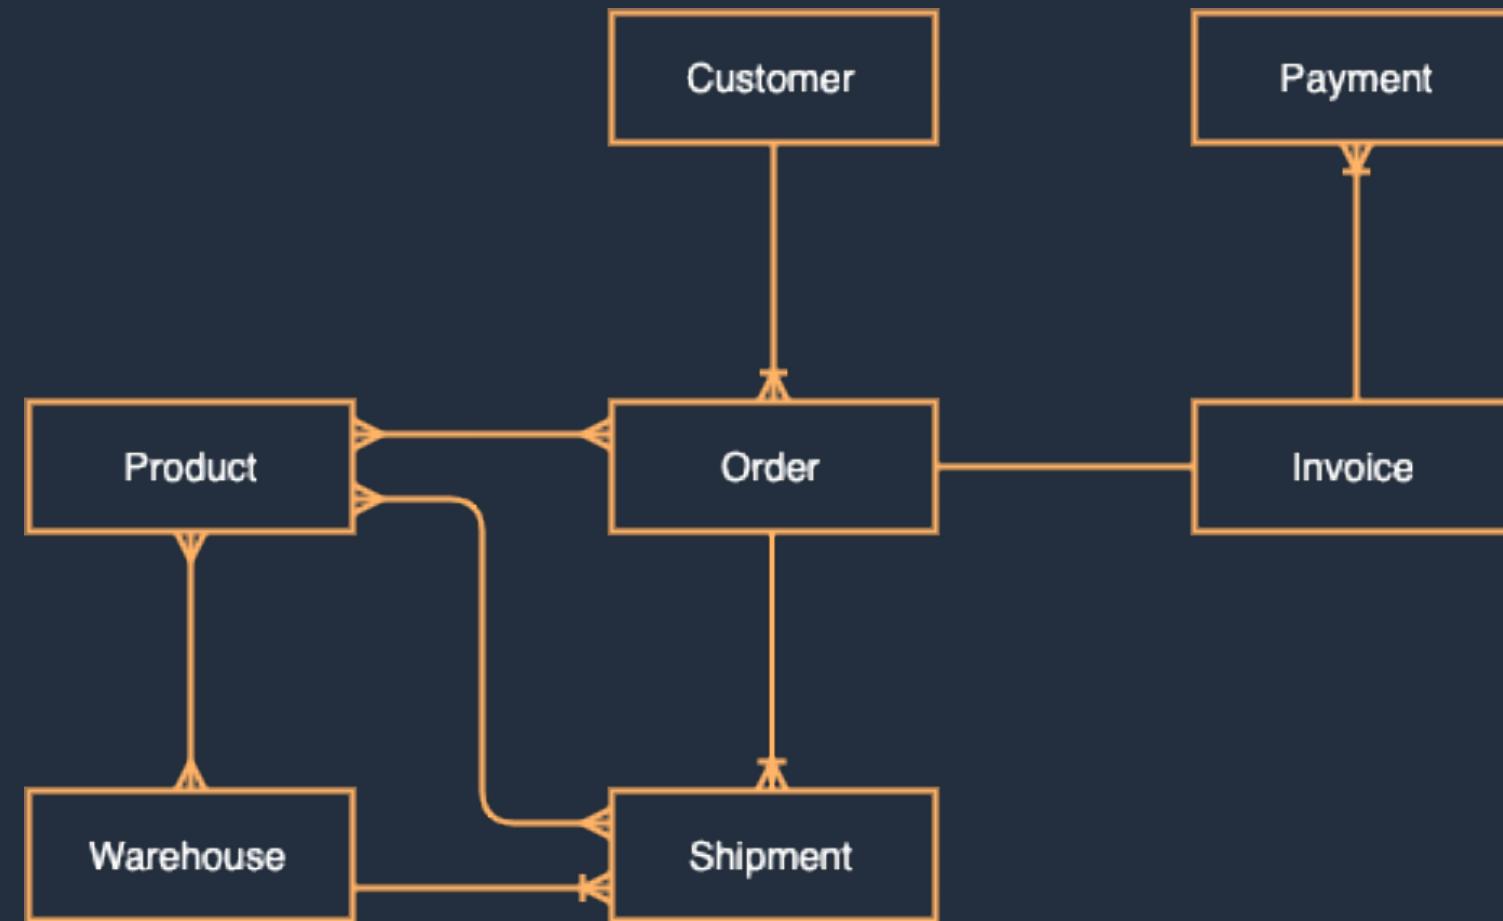
...

A **customer** visits an online shop, browses through different **products** and places an **order** for some of the products. Based on the **invoice**, customer can **pay** using discount code or gift card and pay for the remaining amount by credit card.

Purchased products will be picked from one or several **warehouses** and will be **shipped** to the provided address.

...

Entity Relation Diagram



Identified Access Patterns

Get customer for a given customerId

Get product for a given productId

Get warehouse for a given warehouseId

Get a product inventory for all warehouses by a productId

Get order for a given orderId

Get all products for a given orderId

Get invoice for a given orderId

Get all shipments for a given orderId

Get all orders for a given productId for a given date range

Get invoice for a given invoiceId

Get all payments for a given invoiceId

Get shipment detail for a given shipmentId

Get all shipments for a given warehouseId

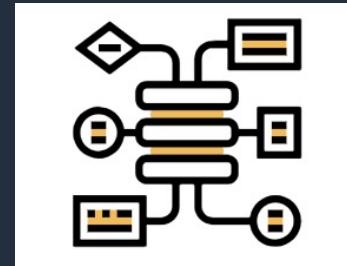
Get inventory of all products for a given warehouseId

Get all invoices for a given customerId for a given date range

Get all products ordered by a given customerId for a given date range

Access Patterns	Table/GSI/LSI	Key Condition	Filter Expression
Get customer for a given customerId			
Get product for a given productId			
Get warehouse for a given warehouseId			
Get a product inventory for all warehouses by a productId			
Get order for a given orderId			
Get all products for a given orderId			
Get invoice for a given orderId			
Get all shipments for a given orderId			
Get all orders for a given productId for a given date range			
Get invoice for a given invoiceId			
Get all payments for a given invoiceId			
Get shipment detail for a given shipmentId			
Get all shipments for a given warehouseId			
Get inventory of all products for a given warehouseId			
Get all invoices for a given customerId for a given date range			
Get all products ordered by a given customerId for a given date range			

NoSQL Workbench



Data modeler



Visualizer



Operation builder

A **client-side application** that helps you build scalable, high-performance data models

Simplifies query development and testing

A rich GUI-based tool that **helps you visualize data models and perform DynamoDB operations**

Available for **Windows, macOS, and Linux**

NoSQL Workbench

To download NoSQL Workbench:



Step-by-step data modeling examples:



<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/workbench.html>

<https://github.com/aws-samples/amazon-dynamodb-design-patterns>

1:M Order to Order Item

Primary key		Attributes			
Partition key: PK	Sort key: SK	EntityType	Date	Price	Quantity
o#12345	c#12345	order	2021-04-20T09:10:00Z		
		orderItem	p#12345	100	2
	p#99887	orderItem	p#99887	40	5
		orderItem	p#12345	100	2

1:M OrderItem to Order

Primary key		Attributes			
Partition key: GSI1-PK	Sort key: GSI1-SK	PK	SK	EntityType	Price
o#54321	c#12345	order	2021-01-19T09:18:00Z		
		orderItem	p#12345	100	2
	p#34567	orderItem	p#34567	50	4
		orderItem	p#12345	100	2

PK	SK	EntityType	Price	Quantity
o#54321	p#12345	orderItem	100	2
o#12345	p#12345	orderItem	100	2
o#54321	p#99887	orderItem	40	5
o#54321	p#34567	orderItem	50	4

Example – Device State Log

High cardinality

Sorted by date time

Filter by Status

Primary key		
Partition key: DeviceID	Sort key: Date	State
d#12345	2020-04-24T14:40:00	WARNING1
d#12345	2020-04-24T14:45:00	WARNING1
d#12345	2020-04-24T14:50:00	WARNING1
d#12345	2020-04-24T14:55:00	WARNING1
d#12345	2020-04-11T05:50:00	NORMAL
d#54321	2020-04-11T05:55:00	WARNING3
d#54321	2020-04-11T06:00:00	WARNING3
d#54321		NORMAL

Access Pattern: Get all logs for a specific device state showing the most recent logs first

```
SELECT * FROM DeviceLog  
WHERE DeviceID = 'd#12345'  
ORDER BY Date DESC  
FILTER ON State='WARNING1'
```

Returned

Primary key		
Partition key: DeviceID	Sort key: Date	
d#12345	2020-04-24T14:40:00	State WARNING1
d#12345	2020-04-24T14:45:00	State WARNING1
d#12345	2020-04-24T14:50:00	State WARNING1
d#12345	2020-04-24T14:55:00	State NORMAL
d#54321	2020-04-11T05:50:00	State WARNING3
d#54321	2020-04-11T05:55:00	State WARNING3
d#54321	2020-04-11T06:00:00	State NORMAL

Filtered

```
aws dynamodb query  
--table-name DeviceStateLog  
--key-condition-expression "#dID = :dID"  
--no-scan-index-forward  
--filter-expression "#s = :s"  
--return-consumed-capacity TOTAL  
--expression-attribute-names '{"#dID":"DeviceID","#s":"State"}'  
--expression-attribute-values '{":dID":{"S":"d#12345"},":s":{"S":"WARNING1"}}'
```

Use Composite Sort Key Instead

Partition key: DeviceID	Primary key	Sort key: State#Date
d#12345		NORMAL#2020-04-24T14:55:00
		WARNING1#2020-04-24T14:40:00
		WARNING1#2020-04-24T14:45:00
		WARNING1#2020-04-24T14:50:00
d#54321		NORMAL#2020-04-11T06:00:00
		NORMAL#2020-04-11T09:30:00
		WARNING2#2020-04-11T09:25:00
		WARNING3#2020-04-11T05:50:00
		WARNING3#2020-04-11T05:55:00

```
aws dynamodb query  
--table-name DeviceStateLog  
--no-scan-index-forward  
--key-condition-expression "#dID = :dID AND begins_with(#s, :sd)"  
--expression-attribute-names '{"#dID":"DeviceID","#s":"State#Date"}'  
--expression-attribute-values '{":dID":{"S":"d#12345"},":sd":{"S":"WARNING1#"} }'  
--return-consumed-capacity TOTAL
```

Access Pattern: Get all device logs for a given operator between two dates

Base Table

Primary key		
Partition key: DeviceID	Sort key: State#Date	Attributes
d#12345	NORMAL#2020-04-24T14:55:00	Operator Liz Date 2020-04-24
	WARNING1#2020-04-24T14:45:00	Operator Liz Date 2020-04-24
	WARNING1#2020-04-24T14:50:00	Operator Liz Date 2020-04-24
	NORMAL#2020-04-11T06:00:00	Operator Liz Date 2020-04-11
	NORMAL#2020-04-11T09:30:00	Operator Sue Date 2020-04-11
	WARNING2#2020-04-11T09:25:00	Operator Sue Date 2020-04-11
d#11223	WARNING3#2020-04-11T05:55:00	Operator Liz Date 2020-04-11
	WARNING4#2020-04-27T16:10:00	Operator Sue Date 2020-04-27
	WARNING4#2020-04-27T16:15:00	Operator Sue Date 2020-04-27

Primary key		Attributes	State#Date	DeviceID
Partition key: Operator	Sort key: Date	State#Date	DeviceID	
Liz	2020-04-11	WARNING3#2020-04-11T05:55:00	d#54321	
	2020-04-11	State#Date	DeviceID	
Sue	2020-04-24	NORMAL#2020-04-11T06:00:00	d#54321	
	2020-04-24	State#Date	DeviceID	
Sue	2020-04-24	WARNING1#2020-04-24T14:45:00	d#12345	
	2020-04-24	State#Date	DeviceID	
Sue	2020-04-24	WARNING1#2020-04-24T14:50:00	d#12345	
	2020-04-24	State#Date	DeviceID	
Sue	2020-04-11	NORMAL#2020-04-24T14:55:00	d#12345	
	2020-04-11	State#Date	DeviceID	
Sue	2020-04-11	WARNING2#2020-04-11T09:25:00	d#54321	
	2020-04-11	State#Date	DeviceID	
Sue	2020-04-11	NORMAL#2020-04-11T09:30:00	d#54321	
	2020-04-11	State#Date	DeviceID	
Sue	2020-04-27	WARNING4#2020-04-27T16:10:00	d#11223	
	2020-04-27	State#Date	DeviceID	
Sue	2020-04-27	WARNING4#2020-04-27T16:15:00	d#11223	
	2020-04-27	State#Date	DeviceID	

GSI-Operator

Access Pattern: Get all device logs for a given operator between two dates

Primary key		Attributes	
Partition key: Operator	Sort key: Date		
Liz	2020-04-11	State#Date	DeviceID
		WARNING3#2020-04-11T05:55:00	d#54321
	2020-04-11	State#Date	DeviceID
		NORMAL#2020-04-11T06:00:00	d#54321
	2020-04-24	State#Date	DeviceID
		WARNING1#2020-04-24T14:45:00	d#12345
	2020-04-24	State#Date	DeviceID
		WARNING1#2020-04-24T14:50:00	d#12345
	2020-04-24	State#Date	DeviceID
		NORMAL#2020-04-24T14:55:00	d#12345
Sue	2020-04-11	State#Date	DeviceID
		WARNING2#2020-04-11T09:25:00	d#54321
	2020-04-11	State#Date	DeviceID
		NORMAL#2020-04-11T09:30:00	d#54321
	2020-04-27	State#Date	DeviceID
		WARNING4#2020-04-27T16:10:00	d#11223
	2020-04-27	State#Date	DeviceID
		WARNING4#2020-04-27T16:15:00	d#11223

```
aws dynamodb query  
--table-name DeviceStateLog  
--index-name GSI1  
--key-condition-expression "#op = :op AND #d between :d1 AND :d2"  
--expression-attribute-names '{"#op": "Operator", "#d": "Date"}'  
--expression-attribute-values '{":op": {"S": "Liz"}, ":d1": {"S": "2020-04-20"}, ":d2": {"S": "2020-04-25"}}'  
--return-consumed-capacity TOTAL
```

GSI – Operator

Access Pattern: Get all escalated logs for a given supervisor

GSI-Supervisor

Base Table

Primary key		Attributes	
Partition key: DeviceID	Sort key: State#Date	Operator	Date
d#12345	NORMAL#2020-04-24T14:55:00	Liz	2020-04-24
	WARNING1#2020-04-24T14:45:00	Liz	2020-04-24
	WARNING1#2020-04-24T14:50:00	Liz	2020-04-24
d#54321	NORMAL#2020-04-11T06:00:00	Liz	2020-04-11
	NORMAL#2020-04-11T09:30:00	Sue	2020-04-11
	WARNING2#2020-04-11T09:25:00	Sue	2020-04-11
d#11223	WARNING3#2020-04-11T05:55:00	Liz	2020-04-11
	WARNING4#2020-04-27T16:10:00	Sue	2020-04-27
	WARNING4#2020-04-27T16:15:00	Sue	2020-04-27

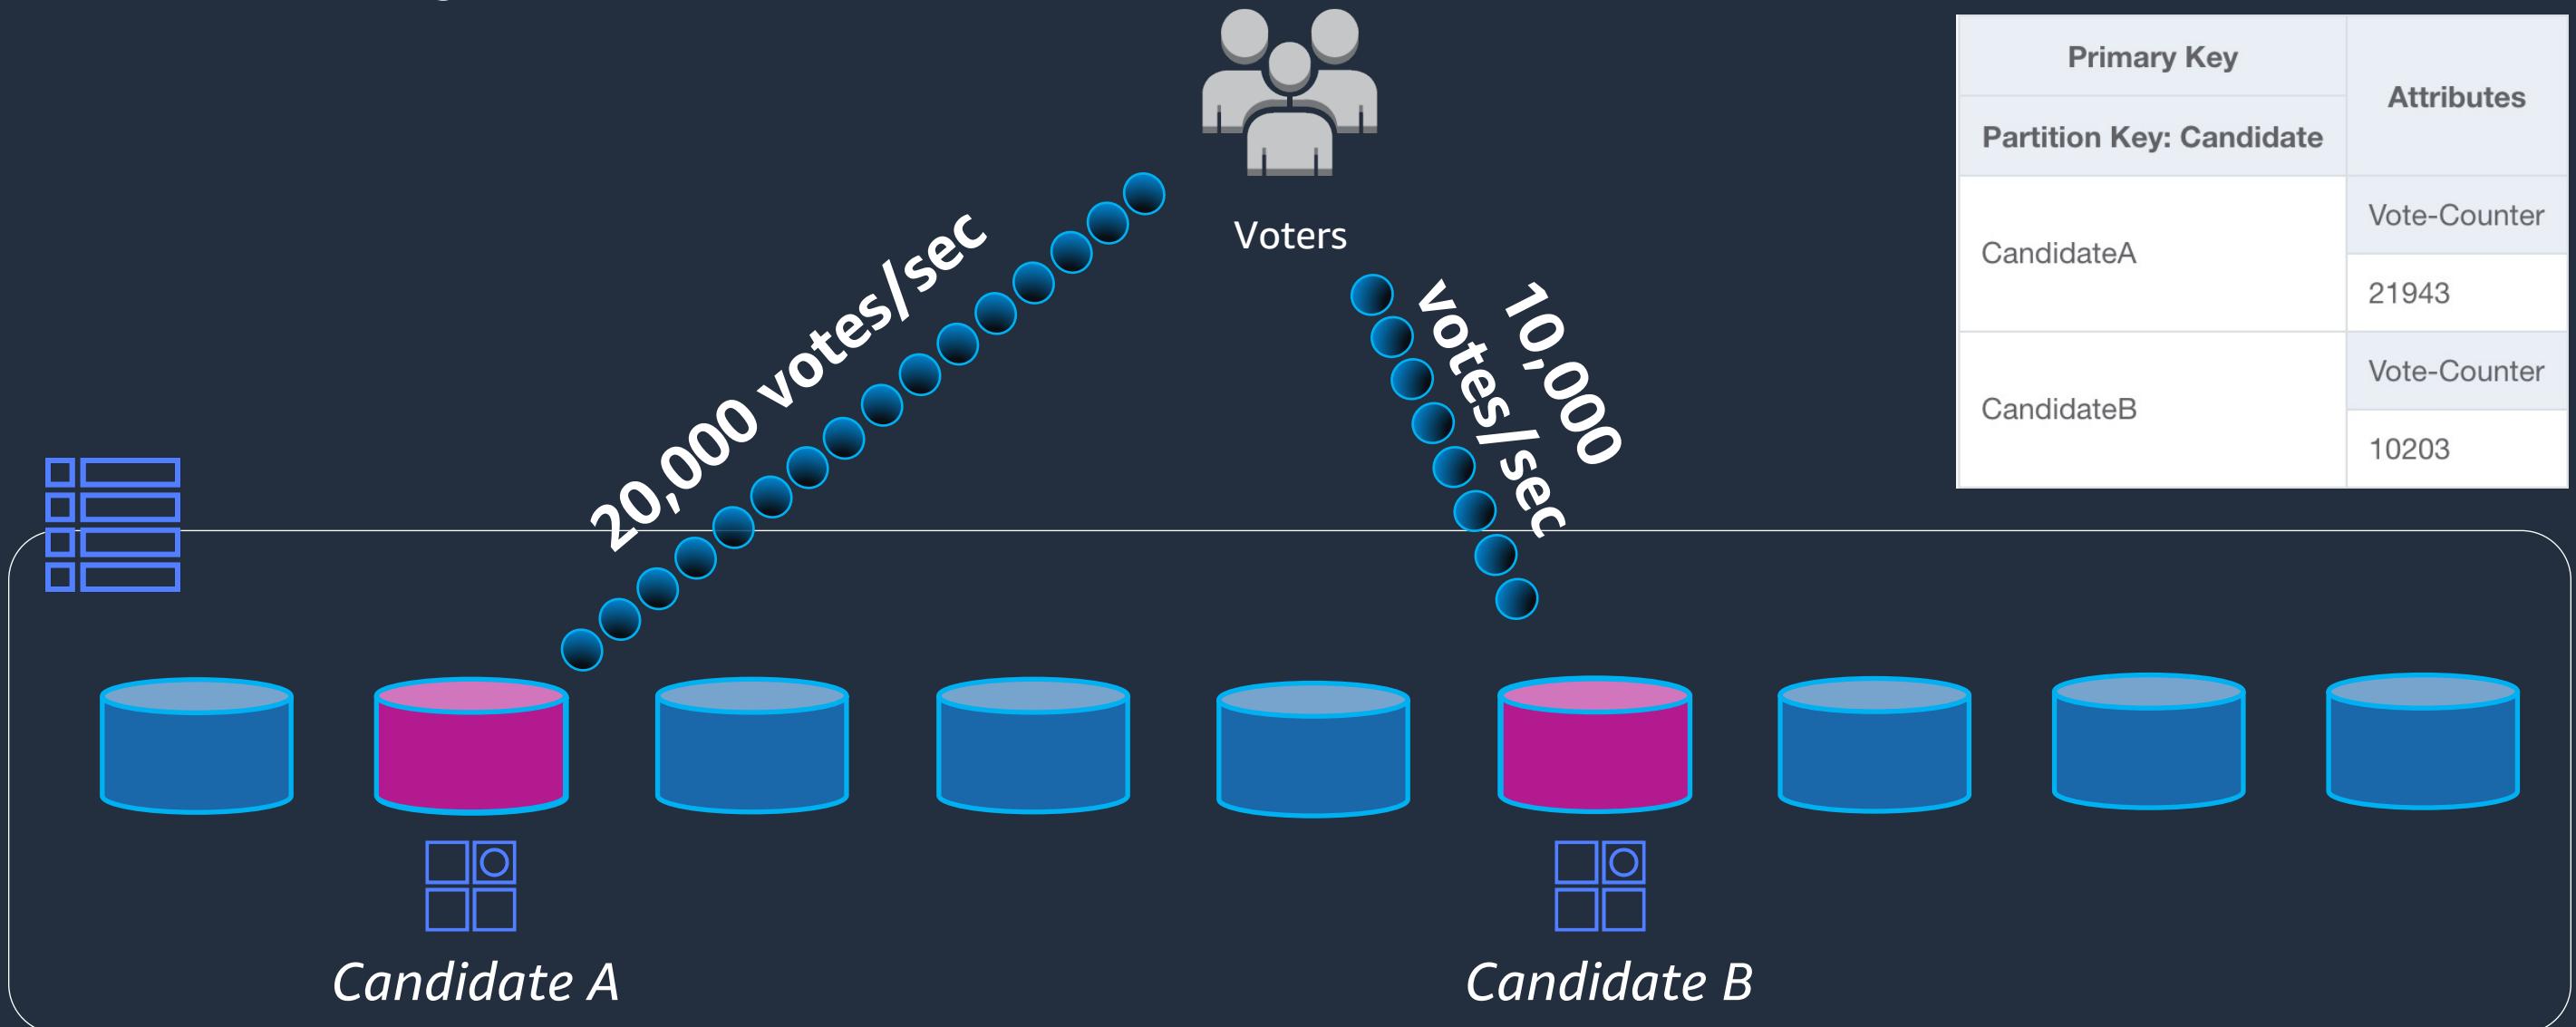
Primary key		Attributes	
Partition key: EscalatedTo	Sort key: State#Date	DeviceID	Operator
Sara	WARNING4#2020-04-27T16:15:00	d#11223	Sue

Sparse GSI: Only items that match the GSI index are projected.

Good for:

- ‘Needle in the haystack’
- Cost effective ‘scans’
- Item management

Example – Voting: Scaling high write throughput & low cardinality...



Single Item limit: 1000 WCUs

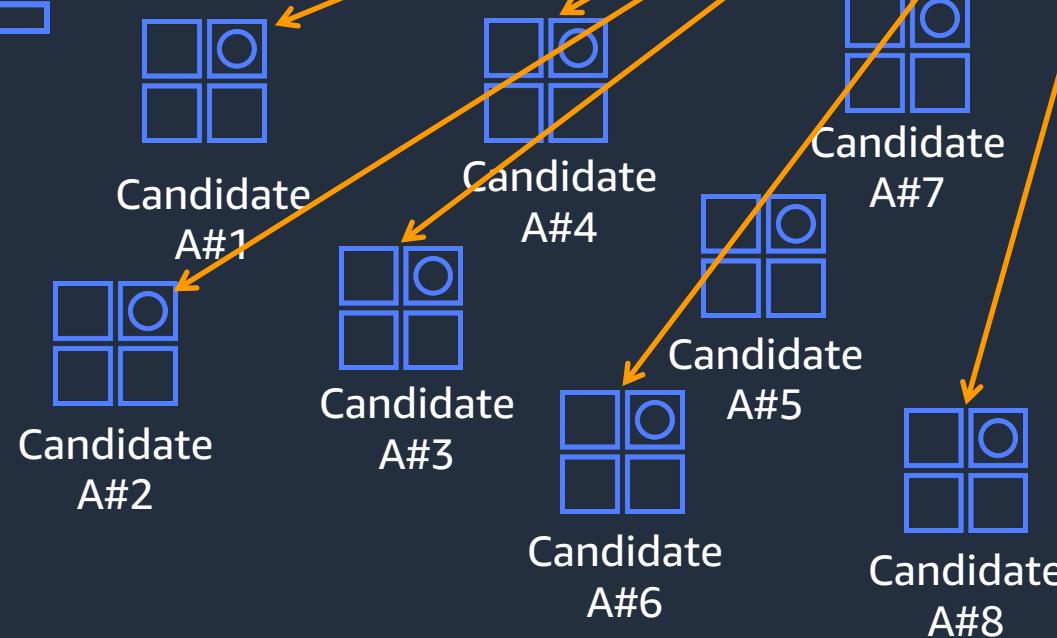
Write shard the partition key...

UPDATE Item :

{

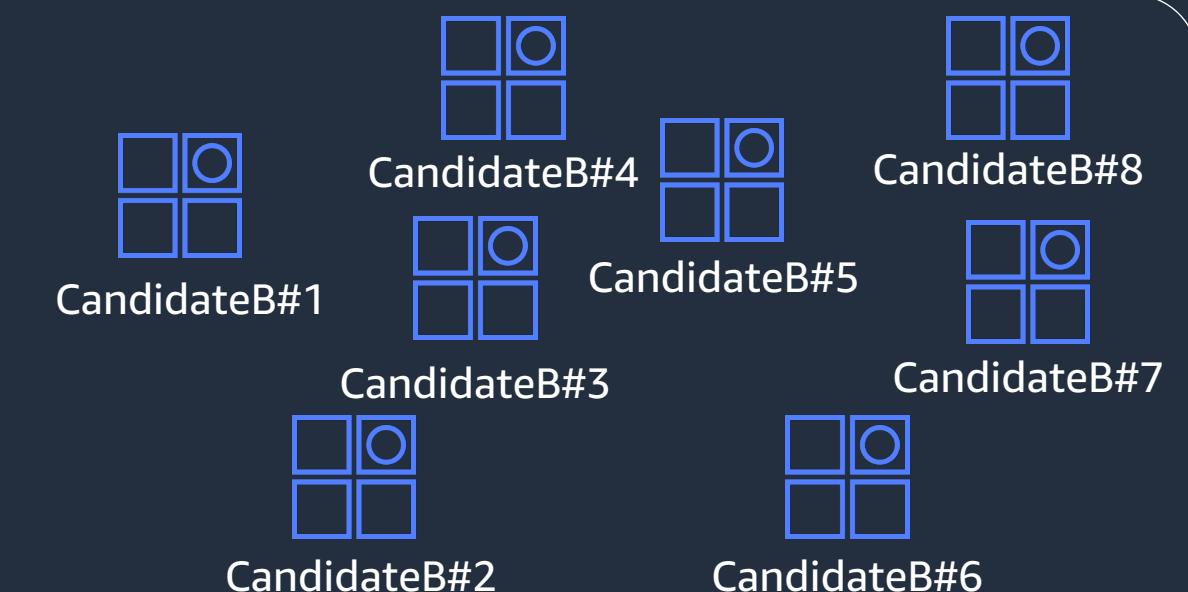
**“CandidateA#”+rand(0,N),
Vote-Counter +1**

}

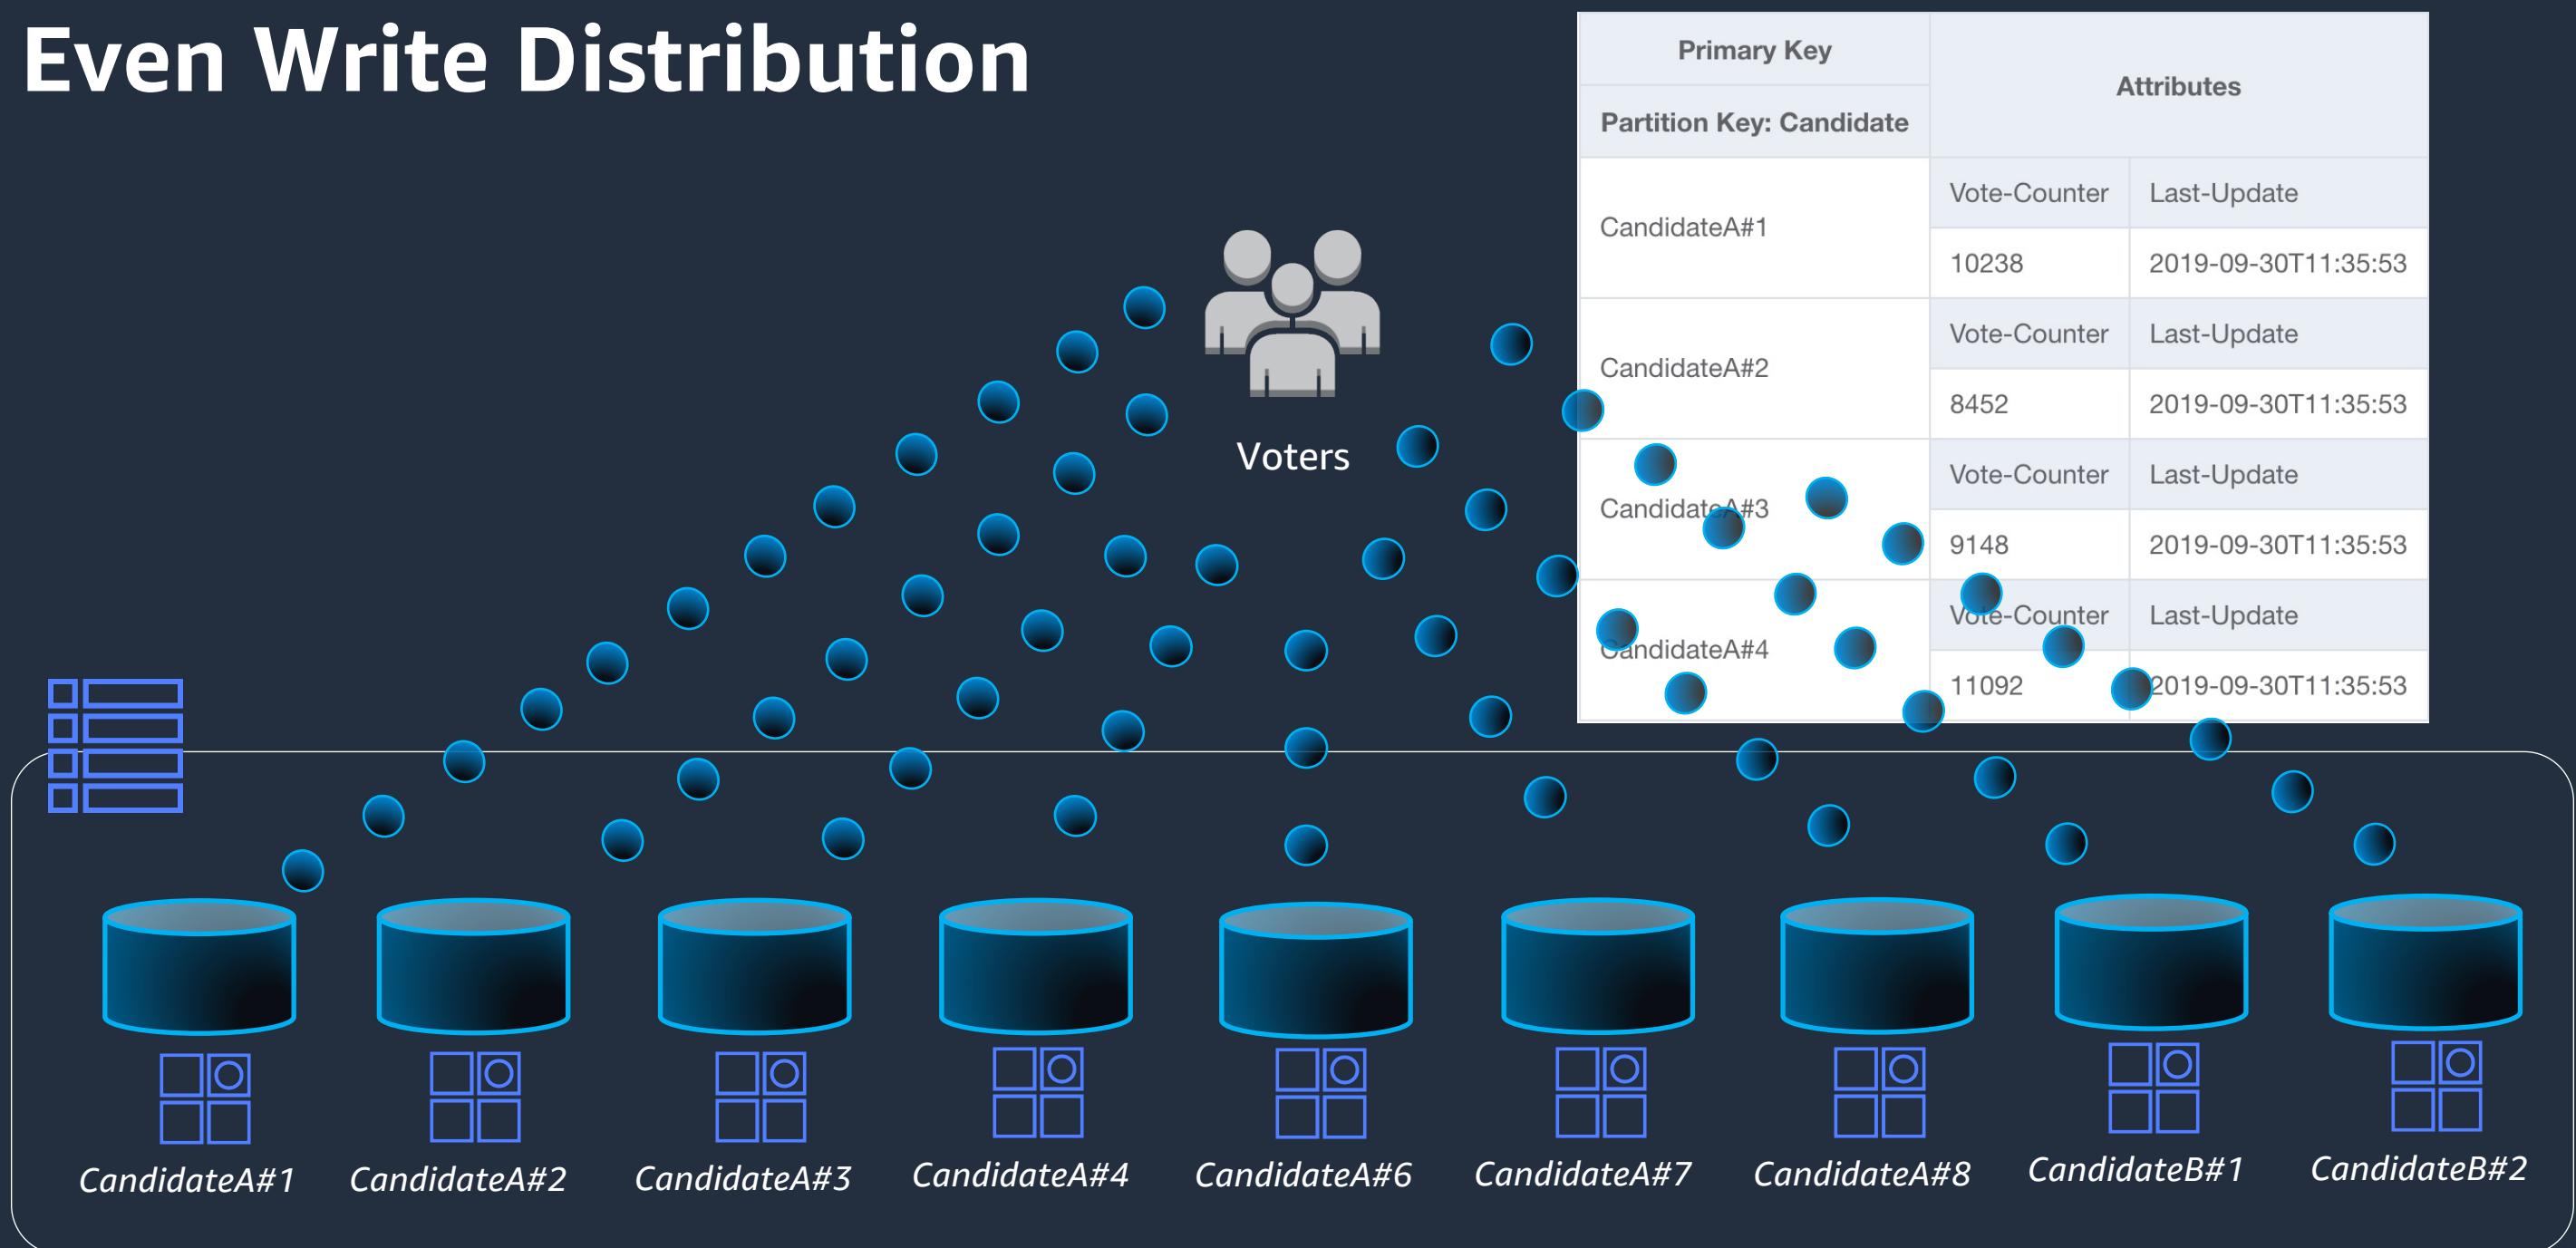


Voters

Primary Key	Attributes	
	Partition Key: Candidate	
CandidateA#1	Vote-Counter	Last-Update
	10238	2019-09-30T11:35:53
CandidateA#2	Vote-Counter	Last-Update
	8452	2019-09-30T11:35:53
CandidateA#3	Vote-Counter	Last-Update
	9148	2019-09-30T11:35:53
CandidateA#4	Vote-Counter	Last-Update
	11092	2019-09-30T11:35:53



Even Write Distribution



$$\text{Item Throughput} = N * 1000 \text{ WCUs}$$

How many artificial shards? (*writes*)

$$\frac{100000 \text{ Votes per second} * 1 \text{ WCU} / \text{Each Vote item consumes}}{1000 \text{ Total WCU per partition}} = 100 \text{ Artificial shards}$$

1 KB * (1 WCU / 1 KB) * (1) = 1 WCU

↑ ↑ ↑

Vote item size Conversion ratio Standard writes

↑ ↑ ↑

Votes per second Each Vote item consumes Total WCU per partition Artificial shards

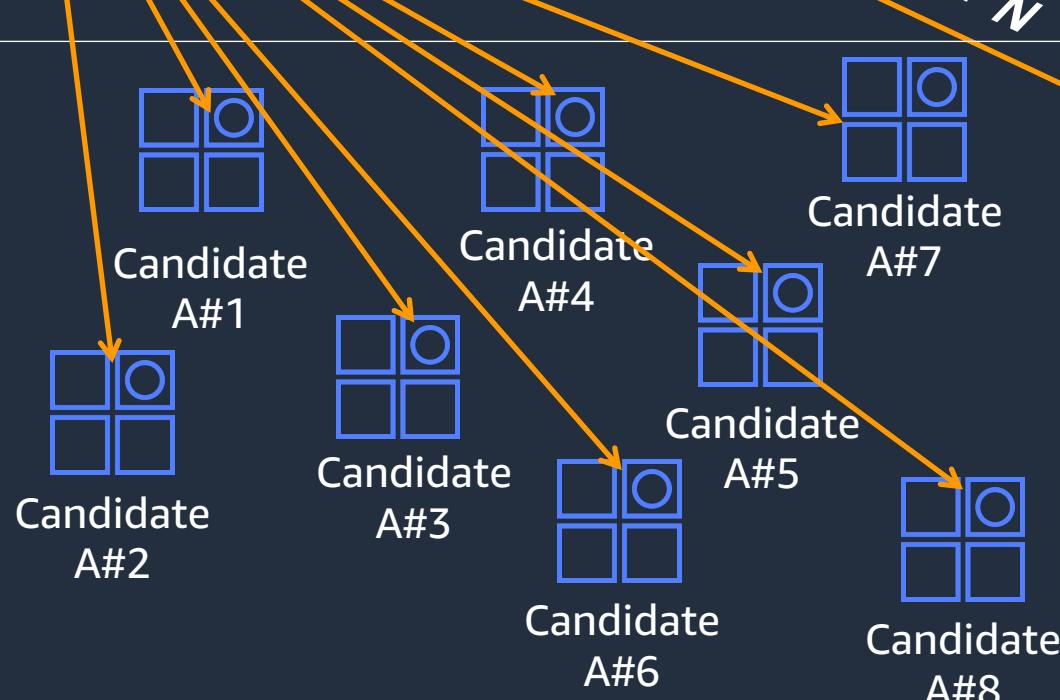


Retrieve result: (in parallel)



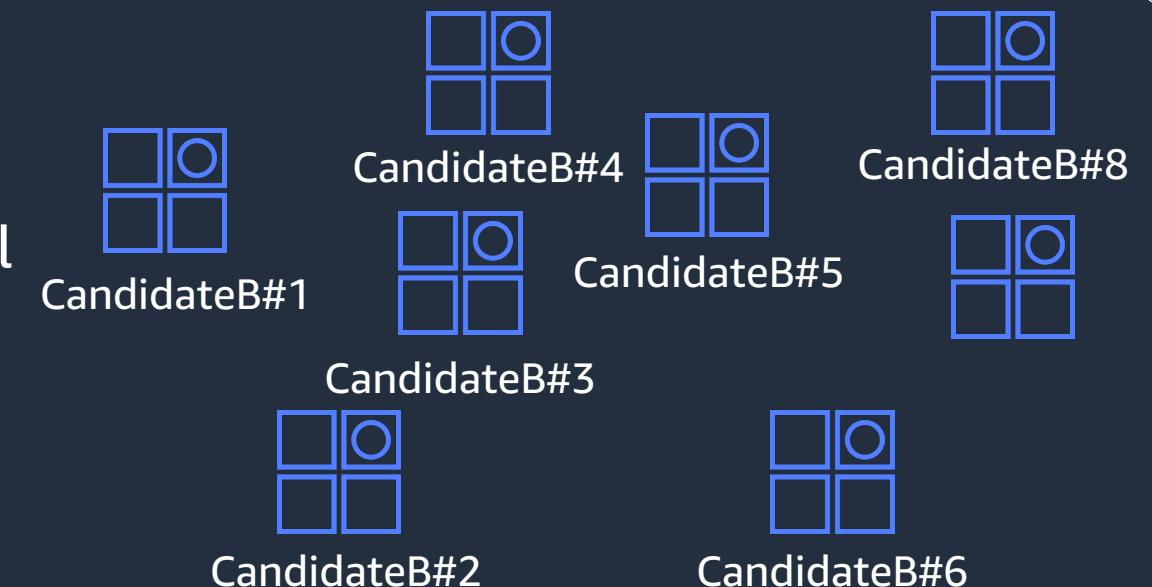
1. Collect

2. Store



3. Retrieve

Primary Key	Attributes		
	Vote-Counter	Last-Update	Total
CandidateA#Total	28692	2019-09-30T11:35:53	28692
CandidateA#2	8452	2019-09-30T11:35:53	
CandidateA#3	9148	2019-09-30T11:35:53	
CandidateA#4	11092	2019-09-30T11:35:53	



Example – Shopping Cart : Document Indexing

```
{  
    "UserProfile" : {  
        "FirstName": "Paul",  
        "LastName": "Atreides",  
        "DateJoined": "1965-08-01"},  
    "Store" : {  
        "store_id": "STOREUID",  
        "city": "Los Angeles",  
        "zip_code": "90029"}  
    "ShoppingCart" : [  
        {"Spice":  
            { "SKU": "SpiceSKU",  
                "CategoryID": "FictionalSpice",  
                "DateAdded": "2019-06-11"},  
            {"EspressoBeans":  
                { "SKU": "CaffeineSKU",  
                    "CategoryID": "FOODANDDRINK",  
                    "DateAdded": "2019-06-10"}}],  
    "shippingAddress" : {  
        "street_address": "1234 Arrakis Dr",  
        "city": "Los Angeles",  
        "zip_code": "90029",  
        "status": "default"}  
    "OrderHistory#OrderUID" : {  
        "ProductA": "SKU_A",  
        "ProductB": "SKU_B",  
        "DateOrdered": "2018-09-28"}  
}
```

Vertical partitioning

Primary key		Attributes		
Partition key: PK	Sort key: SK	Selectively query 'nested' attributes		
UserID	Address#USA#CA#LA#90029	data	GSI-SK	
		"Street Address"	Default	
		data	GSI-SK	
		CoffeeSKU	2019-11-27T103324	
		data	GSI-SK	
		SpiceSKU	2019-11-28T091245	
		data	GSI-SK	
Cart#ACTIVE#Coffee		BEGINS_WITH 'Cart#ACTIVE'		
Cart#ACTIVE#Spice				
Cart#SAVED#Cocoa				
OrderHistory#OrderUID				
ProfileName				
Store#StoreUID				

Fetch items for a specific user that are active in the shopping cart.

BEGINS_WITH 'Cart#ACTIVE'

- Optimize object size
- Selective Queries
- Reduce capacity and cost
- Improve App performance
- Overload GSI on SK...

Takeaways

- Build targeted Queries using composite keys
- Start with Sort Key conditions
- Supplement with Filter Expressions
- Use Global Secondary Index to create new item collections

Find these examples at:

<https://github.com/aws-samples/amazon-dynamodb-design-patterns>

Please complete the survey!

<https://eventbox.dev/survey/XHCETBL>

We value your feedback.



Thank you!

Visit DynamoDB.com

Follow [@DynamoDB](https://twitter.com/DynamoDB) on 

